

Algorithmic aspects of vertex elimination on graphs

Advanced Algorithms project - A.Y. 2020/21

July 23, 2021


Marco Bonelli
10522665

Data structures

I used the [Boost Graph Library](#) to create a generic templated implementation of the algorithms that can work on different implementations of graphs.

Graphs: `boost::adjacency_list<vecS, vecS, undirectedS>`

This is the Graph implementation that I used, a vector of vectors:
fast iteration over vertices and neighbors of a vertex.




Vertex orderings: `std::vector<Graph::vertex_descriptor>`

Vector of vertex descriptors, which for `adjacency_list<vecS, vecS, ...>` simply means `size_t`.

FILL algorithm

Computes the **chordal completion** of any simple, connected, undirected graph given an elimination order for its vertices. Expected complexity: $O(V+E)$ time and space.

This means modifying the graph adding all necessary edges belonging to the *fill-in* of the graph. The *fill-in* is the union of the *deficiencies* of each vertex following the elimination order.



FILL algorithm

Define a mapping $v \rightarrow$ set of successors of v . The successors of a vertex \mathbf{v} are the vertices adjacent to \mathbf{v} that come after \mathbf{v} in the order.

- For each vertex \mathbf{v} in order:
 - Find \mathbf{m} as the successor of \mathbf{v} with minimum index in the order
 - For each other successor \mathbf{w} of \mathbf{v} that is not a successor of \mathbf{m} :
 - Connect \mathbf{w} and \mathbf{m} adding the edge (\mathbf{w}, \mathbf{m}) to the graph
 - Add \mathbf{w} to the successors of \mathbf{m}

FILL algorithm


Used data structures:

- `unordered_map<Vertex, Index>` to keep track of the index in the order of each vertex
- `unordered_map<Vertex, unordered_set<Vertex>>` to keep track of the successors of each vertex (this is not necessarily needed but avoids checking all neighbors of each vertex every time)

LEX P algorithm

Computes a **perfect vertex elimination order** for a perfect elimination graph (i.e. a graph for which the existence of such an order is guaranteed, also known as *chordal* graph). Expected complexity: $O(V+E)$ time and space.

A perfect elimination order results in an empty fill-in. In other words FILL would add no edges to the graph given this order.



LEX P algorithm

Create a linked list of labels (initially only one label), each label keeps track of which vertices have the label through a set. The head of the linked list has the highest label.

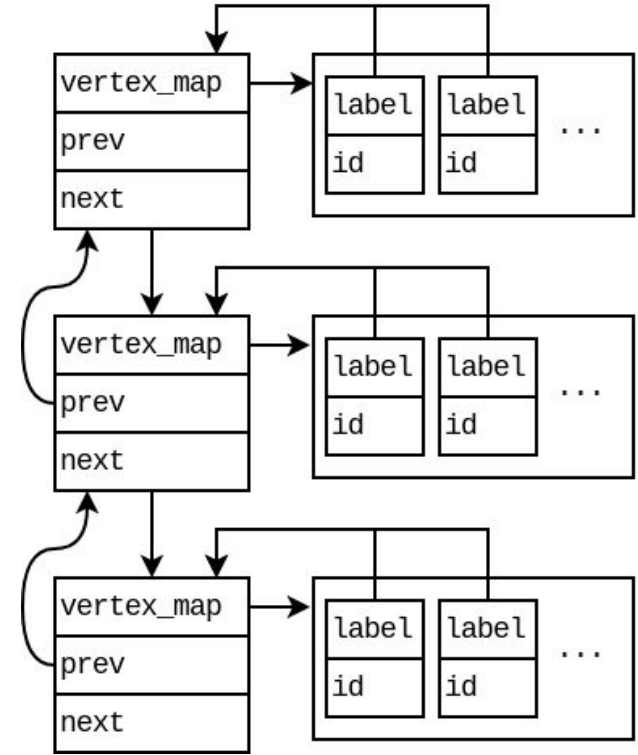
Repeat as many times as the number of vertices:

- Get v as any of the highest-labeled unnumbered vertices, add it as last in the order and mark it as numbered
- Increase the label of each unnumbered neighbor of v by creating a new label and inserting it right before the previous label. Vertices with the same old label need to end up with the same new label (use a map old \rightarrow new to keep track of this).

LEX P algorithm

Used data structures:

- `struct Label {vertex_map, prev, next}`
- `struct LabeledVertex {label, id}`
- `unordered_map<Vertex, LabeledVertex*>`
for unnumbered vertices
- `unordered_map<Label*, Label*>`
to temporary hold newly created labels before
insertion



LEX M algorithm

Computes a **minimal vertex elimination order** of any simple, connected, undirected graph. Expected complexity: $O(VE)$ time, $O(V+E)$ space.

A minimal elimination order is such if there is no other elimination order that produces a fill-in that is a strict subset of the one produced by it.

LEX M algorithm

Repeat as many times as the number of vertices:

1. Get \mathbf{v} as the highest-labeled unnumbered vertex, add it as last in the order and mark it as numbered
2. Increase the label of every unnumbered vertex \mathbf{w} that is reachable from \mathbf{v} *through a chain of unnumbered lower-labeled (than \mathbf{w}) vertices.*

In step 2, the unnumbered vertices are explored in lexicographical breadth-first order of label by keeping one queue of vertices for each different label. Initially queues only hold the neighbors of \mathbf{v} , and are then processed until empty in increasing label order.

LEX M algorithm

Labels are simply even unsigned integers from 0 to $2 * n_unique_labels - 1$. Increasing a label (done at most once) means +1. To keep unnumbered vertices sorted according to their label a radix sort algorithm (LSD base 16) is used at the end of each step of the main loop, and then vertices are re-labeled to even numbers again.

Used data structures:

- `unordered_set<Vertex>` for unnumbered vertices and for reached vertices during lexicographic BFS
- `unordered_map<Vertex, Label>` to keep track of the label of each vertex
- `unordered_map<Label, std::deque<Vertex>>` queues for lexicographic BFS

LEX M algorithm

Errors/tipos in the implementation of the paper:

- The loop should run in the opposite direction with j from 1 to k : the graph needs to be explored starting from lower labeled vertices.
- The $if\ l(z) < k$ should be $if\ l(z) > j$: the label of vertex z should only be increased if we can reach it through a chain of lower-labeled vertices.
- All the k s that appear inside the loop should be j s.

```
search: for  $j:=k$  step  $-1$  until  $1$  do  
  while  $reach(k) \neq \emptyset$  do begin  
    delete a vertex  $w$  from  $reach(k)$ ;  
    for  $z \in adj(w)$  and  $z$  unreached do begin  
      mark  $z$  reached;  
      if  $l(z) < k$  then begin  
        add  $z$  to  $reach(l(z))$ ;  
         $l(z) := l(z) + 1/2$ ;  
        mark  $\{v, z\}$  as an edge of  $G_\alpha^*$ ;  
      end else add  $z$  to  $reach(k)$ ;  
    end end;
```

Random graph generation

For testing and benchmarking purposes we need to generate:

Random connected graphs (FILL, LEX M)

Use an [Erdős–Rényi model](#) to generate a random graph with a given number of vertices and a given probability of having edges between any possible pair of vertices.

Then, identify the connected components of the graph with the help of a union-find and add one edge at a time between two of them until the graph is completely connected.

Random chordal graphs (LEX P)

Implemented as described by the clique-tree algorithm in [Two methods for the generation of chordal graphs \(Markenzon, Vernet, Araujo, 2008\)](#).

This algorithm generates a chordal graph that has a given fixed number of vertices and respects a given upper bound on the number of edges.

Testing

I wrote unit tests using the [Boost Test](#) library, [gcov](#) + [Codecov](#) for code coverage (continuous-integration through [GitHub actions](#)).

1. Random graph generation utilities: ensure graphs are simple, connected or chordal as needed, not complete unless needed.
2. Radix sort correctly sorts
3. **FILL**: chordal completion on known small graphs, empty fill-in for complete graphs given any elimination order, empty fill-in for chordal graphs given a perfect elimination order.
4. **LEX M**: ensure computed elimination order is minimal for connected graphs and perfect for chordal graphs.
5. **LEX P**: ensure computed elimination order is perfect for chordal graphs.

Testing

```
test/unit/test_lex_p.cc(11): Leaving test suite "LexP"; testing time: 256834us
test/unit/test_radix_sort.cc(10): Entering test suite "RadixSort"
test/unit/test_radix_sort.cc(22): Entering test case "radix_sort_works<unsigned char>"
test/unit/test_radix_sort.cc(22): Leaving test case "radix_sort_works<unsigned char>"; testing time: 2386us
test/unit/test_radix_sort.cc(22): Entering test case "radix_sort_works<unsigned short>"
test/unit/test_radix_sort.cc(22): Leaving test case "radix_sort_works<unsigned short>"; testing time: 2726us
test/unit/test_radix_sort.cc(22): Entering test case "radix_sort_works<unsigned int>"
test/unit/test_radix_sort.cc(22): Leaving test case "radix_sort_works<unsigned int>"; testing time: 3443us
test/unit/test_radix_sort.cc(22): Entering test case "radix_sort_works<unsigned long>"
test/unit/test_radix_sort.cc(22): Leaving test case "radix_sort_works<unsigned long>"; testing time: 4772us
test/unit/test_radix_sort.cc(10): Leaving test suite "RadixSort"; testing time: 13361us
test/unit/test_fill.cc(22): Entering test suite "Fill"
test/unit/test_fill.cc(52): Entering test case "known_graph"
test/unit/test_fill.cc(52): Leaving test case "known_graph"; testing time: 788us
test/unit/test_fill.cc(82): Entering test case "complete_graph_has_empty_fill_in_for_any_order"
test/unit/test_fill.cc(82): Leaving test case "complete_graph_has_empty_fill_in_for_any_order"; testing time: 2387882us
test/unit/test_fill.cc(110): Entering test case "complete_graph_has_empty_fill_in"
test/unit/test_fill.cc(110): Leaving test case "complete_graph_has_empty_fill_in"; testing time: 1056263us
test/unit/test_fill.cc(128): Entering test case "chordal_graph_has_empty_fill_in"
test/unit/test_fill.cc(128): Leaving test case "chordal_graph_has_empty_fill_in"; testing time: 69356us
test/unit/test_fill.cc(22): Leaving test suite "Fill"; testing time: 3514354us
Leaving test module "Main"; testing time: 32541260us
```

Test module "Main" has passed with:

21 test cases out of 21 passed

277273 assertions out of 277273 passed

Timing benchmarks

Done through the [Google Benchmark](#) library, results plotted through Python3 + [matplotlib](#) + [scikit-learn](#).

Benchmarked all algorithms on random connected graphs with different edge densities, linearly increasing the number of vertices. For LEX P: generated a random connected graph, then used LEX M to find a minimal elimination order, then FILL to make it chordal.

Tested edge densities (d): {0.1, 0.25, 0.50, 0.66, 0.75, 1.0}.

Tested number of vertices (V) for each density: {100, 200, ..., 1000}.

Timing benchmarks

```
template <unsigned num, unsigned div>
void lex_m_random_graph(benchmark::State& state) {
    const unsigned v = state.range(0);
    auto g = gen_random_connected_graph<Graph>(v,
(double)num/div);

    for (auto _ : state)
        benchmark::DoNotOptimize(lex_m(g));

    const auto n = boost::num_vertices(g) * boost::num_edges(g);
    state.counters["n"] = n;
    state.counters["v"] = boost::num_vertices(g);
    state.SetComplexityN(n);
}
```

```
BENCHMARK_TEMPLATE(lex_m_random_graph, 1, 10)->DenseRange(100, 1000, 100)->Complexity(benchmark::oN)->Unit(benchmark::kMillisecond);
BENCHMARK_TEMPLATE(lex_m_random_graph, 1, 4)->DenseRange(100, 1000, 100)->Complexity(benchmark::oN)->Unit(benchmark::kMillisecond);
BENCHMARK_TEMPLATE(lex_m_random_graph, 1, 2)->DenseRange(100, 1000, 100)->Complexity(benchmark::oN)->Unit(benchmark::kMillisecond);
BENCHMARK_TEMPLATE(lex_m_random_graph, 2, 3)->DenseRange(100, 1000, 100)->Complexity(benchmark::oN)->Unit(benchmark::kMillisecond);
BENCHMARK_TEMPLATE(lex_m_random_graph, 3, 4)->DenseRange(100, 1000, 100)->Complexity(benchmark::oN)->Unit(benchmark::kMillisecond);
BENCHMARK_TEMPLATE(lex_m_random_graph, 1, 1)->DenseRange(100, 1000, 100)->Complexity(benchmark::oN)->Unit(benchmark::kMillisecond);
```

Timing benchmarks

Benchmark	Time	CPU	Iterations	UserCounters...
fill_in_random_graph<1, 10>/100	0.584 ms	0.583 ms	1266	n=583 v=100
fill_in_random_graph<1, 10>/200	2.78 ms	2.78 ms	250	n=2.214k v=200
fill_in_random_graph<1, 10>/300	7.12 ms	7.12 ms	99	n=4.788k v=300
fill_in_random_graph<1, 10>/400	15.8 ms	15.8 ms	50	n=8.378k v=400
fill_in_random_graph<1, 10>/500	23.8 ms	23.8 ms	30	n=12.81k v=500
fill_in_random_graph<1, 10>/600	41.9 ms	41.9 ms	16	n=18.457k v=600
fill_in_random_graph<1, 10>/700	65.0 ms	65.0 ms	10	n=25.352k v=700
fill_in_random_graph<1, 10>/800	103 ms	103 ms	7	n=33.084k v=800
fill_in_random_graph<1, 10>/900	134 ms	134 ms	5	n=41.238k v=900
fill_in_random_graph<1, 10>/1000	163 ms	163 ms	4	n=51.079k v=1000
fill_in_random_graph<1, 10>_BigO	3030.58 N	3029.99 N		
fill_in_random_graph<1, 10>_RMS	17 %	17 %		
fill_in_random_graph<1, 4>/100	0.728 ms	0.728 ms	929	n=1.346k v=100
fill_in_random_graph<1, 4>/200	3.11 ms	3.11 ms	226	n=5.177k v=200
fill_in_random_graph<1, 4>/300	7.41 ms	7.41 ms	99	n=11.588k v=300
fill_in_random_graph<1, 4>/400	13.8 ms	13.8 ms	47	n=20.601k v=400
fill_in_random_graph<1, 4>/500	24.6 ms	24.6 ms	30	n=31.469k v=500
fill_in_random_graph<1, 4>/600	41.9 ms	41.9 ms	13	n=45.535k v=600
fill_in_random_graph<1, 4>/700	58.0 ms	58.0 ms	12	n=61.964k v=700
fill_in_random_graph<1, 4>/800	77.1 ms	77.1 ms	8	n=80.824k v=800
fill_in_random_graph<1, 4>/900	122 ms	122 ms	6	n=102.209k v=900
fill_in_random_graph<1, 4>/1000	137 ms	137 ms	5	n=126.139k v=1000
fill_in_random_graph<1, 4>_BigO	1056.85 N	1056.70 N		
fill_in_random_graph<1, 4>_RMS	15 %	15 %		
fill_in_random_graph<1, 2>/100	0.732 ms	0.732 ms	887	n=2.56k v=100

Time complexity


Number of edges (E) directly depends on the number of vertices and the edge density (d) (i.e. Erdős–Rényi probability parameter): $E \approx dV(V-1)/2 \approx dV^2$.

Time complexity:


FILL	$O(V + E) = O(V + dV(V - 1)/2) = O(V + V^2) = O(V^2)$
LEX P	$O(V + E) = O(V + dV(V - 1)/2) = O(V + V^2) = O(V^2)$
LEX M	$O(VE) = O(VdV(V - 1)/2) = O(V^3)$

Time complexity: FILL

$T = f(V)$ w/ quadratic regression




$T = f(V + E)$ w/ linear regression




Time complexity: LEX P

$T = f(V)$ w/ quadratic regression




$T = f(V + E)$ w/ linear regression




Time complexity: LEX M

$T = f(V)$ w/ cubic regression



$T = f(V + E)$ w/ linear regression



Memory benchmarks

Done through glibc's `__{malloc,realloc,free}_hook` in a very simple custom memory profiling library to keep track of maximum heap usage.

```
template <unsigned num, unsigned div>
void fill_random_graph(void) {
    char name[128];

    for (const auto v : vertices) {
        Graph g = gen_random_connected_graph<Graph>(v, (double)num/div);
        const unsigned n = boost::num_vertices(g) + boost::num_edges(g);

        sprintf(name, "fill_random_graph<%u,%u> v=%u n=%u", num, div, v, n);
        start_trace(name);
        do_not_optimize(lex_m(g));
        stop_trace();
    }
}
```

```
static void start_trace(const char *name) {
    __malloc_hook      = malloc_hook;
    __realloc_hook      = realloc_hook;
    __free_hook         = free_hook;
    cur_trace_name      = name;
    allocated_memory    = 0;
    max_allocated_memory = 0;
}

static void stop_trace(void) {
    __malloc_hook = NULL;
    __realloc_hook = NULL;
    __free_hook   = NULL;

    std::cout << cur_trace_name
        << ": max " << max_allocated_memory
        << " bytes" << std::endl;
}
```

Space complexity


As in timing benchmarks, the Number of edges is still: $E \approx dV(V-1)/2 \approx dV^2$.

Space complexity (in terms of **additional** space required):


FILL	$O(V + E) = O(V + dV(V - 1)/2) = O(V + V^2) = O(V^2)$
LEX P	$O(V + E) = O(V + dV(V - 1)/2) = O(V + V^2) = O(V^2)$
LEX M	$O(V)$

Space complexity: FILL

$S = f(V)$ w/ quadratic regression




$S = f(V + E)$ w/ linear regression




Space complexity: LEX P

$S = f(V)$ w/ quadratic regression



$S = f(V + E)$ w/ linear regression



Space complexity: LEX M

